



Advanced Database Systems:

Part III.5 Column-Oriented Databases

Reinhard Pichler

Institute of Logic and Computation
"Databases & Artificial Intelligence" Group

Summer Semester 2023

References

❖ Textbooks:

- Pramod J. Sadalage and Martin Fowler: NoSQL Distilled. Addison-Wesley 2013
- Lena Wiese: Advanced Data Management for SQL, NoSQL, Cloud and Distributed Databases, de Gruyter, 2015

❖ Further articles and books:

- Daniel Abadi, Peter Boncz, Stavros Harizopoulos, Stratos Idreos, and Samuel Madden: The Design and Implementation of Modern Column-Oriented Database Systems. Foundations and Trends in Databases, 5(3): 197-280, 2013

❖ Copyright of all figures is with one of these resources.

Contents

❖ Column Stores:

- Basic Ideas
- Explicit vs. Implicit IDs
- Compression
- Late Materialization
- Updates

❖ Column Family Stores (Wide Column Stores):

- Motivation
- Data Model

Contents

❖ Column Stores:

➤ **Basic Ideas**

- Explicit vs. Implicit IDs
- Compression
- Late Materialization
- Updates

❖ Column Family Stores (Wide Column Stores):

- Motivation
- Data Model

Column Stores

❖ Basic Idea:

- Physically store data in **vertical fragments**
- This is in contrast to traditional RDBMSs, which store whole records together (= row store)

Sales				
	saleid	prodid	date	region
1				
2				
3				
4				
5				
6				
7				
8				
9				
10				

(a) Column Store with Virtual Ids

Sales				
	saleid	prodid	date	region
1		1		1
2		2		2
3		3		3
4		4		4
5		5		5
6		6		6
7		7		7
8		8		8
9		9		9
10		10		10

(b) Column Store with Explicit Ids

Sales				
	saleid	prodid	date	region
1				
2				
3				
4				
5				
6				
7				
8				
9				
10				

(c) Row Store

Why Column Stores?

❖ Business Analytics & Data Warehouses:

- For analytics you want to store as much information as possible, hence tables contain many columns.
- Concrete query usually involves aggregation over few columns, e.g., product sales in a given month: involves only product id, date, maybe amount.

❖ With Column Stores:

- Save I/O by reading only the columns you need;
- further optimizations of column-wise processing.

Historical Perspective

❖ Many of the ideas are old:

- Column-oriented layouts were proposed in the 1970s as part of research on time-series databases.
- Why the recent (around 2000) surge in popularity?

❖ Development of server hardware since the 1980s:


- CPUs have become $\sim 10,000$ times faster.
- Cache/RAM/Disk size has also grown by a factor of $\sim 10,000$
- Transfer from disk has become ~ 100 times faster.
- Seek times have become ~ 10 times faster

❖ Consequences:

- Trade-off between CPU and I/O has changed;
- important for column-oriented optimizations.

Vectorized Processing

- Iterating over single columns instead of rows enables further speedups beyond the saved I/O.
- Modern CPUs support **single instruction, multiple data** (SIMD) operations: particularly well suited for column-wise processing
- Even without SIMD: Better cache locality in comparison to iterating over rows.



Source: <https://kernel.org/pub/linux/kernel/people/geoff/cell/ps3-linux-docs/>

Systems

❖ Column-oriented extensions for other systems:

- cstore_fdw for PostgreSQL
- "columnstore indexes" in Microsoft SQL Server
- in-memory column store of Oracle
- SAP HANA stores data in column *and* row format

❖ "Pure" column stores:

- Open source: MonetDB, Apache Kudu (on Hadoop), Druid (Apache)
- Commercial: Vertica (commercial fork of earlier C-Store system), Sybase IQ (an SAP company)

Contents

❖ Column Stores:

- Motivation
- **Explicit vs. Implicit IDs**
- Compression
- Late Materialization
- Updates

❖ Column Family Stores (Wide Column Stores):

- Motivation
- Data Model

Need for Ids

- The system needs to be able to reconstruct the rows.
- Example: Which sales were made on which date?
- This means: We need a way to identify which values in the columns belong to the same row. $\Rightarrow$ Ids

Sales				
	saleid	prodid	date	region
1				
2				
3				
4				
5				
6				
7				
8				
9				
10				

(a) Column Store with Virtual Ids

Sales					
saleid		prodid		date	region
1		1		1	
2		2		2	
3		3		3	
4		4		4	
5		5		5	
6		6		6	
7		7		7	
8		8		8	
9		9		9	
10		10		10	

...

(b) Column Store with Explicit Ids

Explicit Ids

- Associate a tuple identifier with every entry in every column
- Disadvantage: data size increases $\Rightarrow$ more I/O

Sales				
	saleid	prodid	date	region
1				
2				
3				
4				
5				
6				
7				
8				
9				
10				

(a) Column Store with Virtual Ids

Sales				
	saleid	prodid	date	region
1	1	1	1	1
2	2	2	2	2
3	3	3	3	3
4	4	4	4	4
5	5	5	5	5
6	6	6	6	6
7	7	7	7	7
8	8	8	8	8
9	9	9	9	9
10	10	10	10	10

(b) Column Store with Explicit Ids

Virtual Ids

- Avoid Id-column by using the offset as virtual Id.
- Efficient lookup via offset: $\text{startOf}(C) + i * \text{width}(C)$
- Disadvantage: problem with (non-fixed length) compression methods and variable-width columns.

Sales				
	saleid	prodid	date	region
1				
2				
3				
4				
5				
6				
7				
8				
9				
10				

(a) Column Store with Virtual Ids

Sales				
	saleid	prodid	date	region
1		1	1	1
2		2	2	2
3		3	3	3
4		4	4	4
5		5	5	5
6		6	6	6
7		7	7	7
8		8	8	8
9		9	9	9
10		10	10	10

(b) Column Store with Explicit Ids

Contents

❖ Column Stores:

- Motivation
- Explicit vs. Implicit IDs
- **Compression**
- Late Materialization
- Updates

❖ Column Family Stores (Wide Column Stores):

- Motivation
- Data Model

Compression – Motivation

- ❖ Analytical queries usually require a lot of I/O and aggregation (cheap for the CPU)
- ❖ Idea: trade CPU time for I/O by storing compressed data
 - Column-oriented layout allows for specialized choice of compression algorithm depending on e.g. data types.
 - Single columns have similar values (high data locality)
⇒ yet better compression possible.
 - Example: Compressing only product price vs. compressing product price + picture + description + ...
 - Drawback: Indexing and offset addressing may become more difficult.

Run-Length Encoding

- Represent consecutive repetitions of values as pair (value, #repetitions).
- Worse than no compression in case of little repetition (or even unique values)
- Disadvantage: Causes issues with virtual ids.

ReaderID
205
205
205
207
207
207
207
205
205
587
587



ReaderID
(value,start,length)
(205,1,3)
(207,4,4)
(205,8,2)
(587,10,2)


Bit-Vector Encoding

- One bit-string with length of the number of entries in the column for each value.
- 1 in i -th position encodes value (cf. Bitmap Index)
- Only useful with few possible values.
(Exception: further compression of the bit-strings.)

ReaderID	205	207	587
205	1	0	0
205	1	0	0
205	1	0	0
207	0	1	0
207	0	1	0
207	0	1	0
207	0	1	0
205	1	0	0
205	1	0	0
587	0	0	1
587	0	0	1


Dictionary Encoding

- Create dictionary from values and use integer indices into the dictionary instead of the actual values.
- Replaces long values with shorter placeholders (e.g., 4 byte integer instead of a name string)
- Can turn a variable-width column into fixed-width



Null Suppression


- Some applications lead to *sparse* columns, i.e., columns with many NULL values.
- Idea: Store sparse columns without the NULL values.
- Example: Position bit-string encoding below. Other null suppression techniques also exist.



Operations on Compressed Data

- ❖ More intricate compression schemes -- with better compression ratios, e.g. Lempel-Ziv, Huffman coding, etc.
- ❖ Ideal solution: Ultimate performance boost if we can **operate directly on the compressed data!** This saves I/O and eliminates the need for decompression.
- ❖ With the presented compression schemes, many common operations are in fact possible without decompression. Example: **SUM over a run-length encoded column.** Directly add $\langle \text{value} \rangle * \# \text{repetitions}$.
- ❖ In general, requires special consideration by the query execution engine. Results in highly complex systems.

Real World Performance (Space Reduction)



- Column with ~100 million 32bit integers.
- Significant improvements even with very simple methods
- Similar improvements observed in query execution time.

Source: Abadi, Madden, Ferreira. "Integrating compression and execution in column-oriented database systems."

Compression – Conclusion

- ❖ There are many compression strategies to consider.
- ❖ Many allow (some) operations on compressed data.
- ❖ Concrete choice of compression strategy is column- and application-specific.
- ❖ **Example:**
 - "date" column in table of invoices: long sequences of same date expected $\Rightarrow$ run-length encoding
 - "month" of birth dates of users: no regularity expected, but only 12 values $\Rightarrow$ bit-vector encoding
 - address column $\Rightarrow$ dictionary

Contents

❖ Column Stores:

- Motivation
- Explicit vs. Implicit IDs
- Compression
- **Late Materialization**
- Updates

❖ Column Family Stores (Wide Column Stores):

- Motivation
- Data Model


Materialization

- ❖ Queries often access multiple attributes of a table. Thus, the separately stored columns need to be brought together at some point. This is called **materialization**.
- ❖ **Early materialization**: Reconstruct (materialize) the rows right after reading the columns. Simple but we lose most of the column-oriented performance improvements:
 - Materialization may require decompression.
 - No or very limited vectorized processing.
- ❖ Better solution: **late materialization**, i.e., materialize as late as possible.


Late Materialization

- ❖ Operate on individual columns as much as possible to exploit column-oriented optimizations.
- ❖ Combine columns as late as possible.
- ❖ For some aggregation tasks, it may be possible to completely skip materialization.
- ❖ Requires keeping track of intermediate position tables


Late Materialization –Example



Late Materialization –Example (Cont.)



Late Materialization –Example (Cont.)



Contents

❖ Column Stores:

- Motivation
- Explicit vs. Implicit IDs
- Compression
- Late Materialization
- **Updates**

❖ Column Family Stores (Wide Column Stores):

- Motivation
- Data Model

Updating Columns

- ❖ Storing columns in separate files means **multiple writes** (all columns) in case of insertion and deletion.
- ❖ Inserting and deleting in compressed columns can be very expensive and complex. In the worst case, this requires: **decompress, update, recompress**.
- ❖ In consequence, many systems manage a separate **write-store** for recent updates in addition to the **read-store** (= main database).
- ❖ Queries are executed on both the read- and the write-store and the results are merged on-the-fly.
- ❖ Contents of the write-store is periodically propagated to the read-store to keep the write-store small.

Contents

❖ Column Stores:

- Motivation
- Explicit vs. Implicit IDs
- Compression
- Late Materialization
- Updates

❖ Column Family Stores (Wide Column Stores):

- **Motivation**
- Data Model

Pioneering Work at Google: Bigtable

Fay Chang, Jeffrey Dean, Sanjay Ghemawat, Wilson C. Hsieh, Deborah A. Wallach, Michael Burrows, Tushar Chandra, Andrew Fikes, Robert E. Gruber:

Bigtable: A Distributed Storage System for Structured Data. ACM Trans. Comput. Syst. 26(2): 4:1-4:26 (2008)

Also: Best paper at OSDI 2006:

(Symposium on Operating Systems Design and Implementation)

Pioneering Work at Google: Bigtable

❖ Motivation: huge amount of data

- Web pages:
contents, crawl metadata, links, anchors, page rank
- Per-user data:
user preference settings, recent queries/search results
- Geographic locations:
physical entities (shops, restaurants, etc.), roads,
satellite image data, user annotations

❖ Principal challenge: scale

- #web pages, many versions per page
- #users, #queries per second
- size of image data

Requirements

❖ Up-to-dateness:

- Data has to be continuously updated by asynchronous processes.
- Allow access to most current data at any time.

❖ Performance requirements:

- Very high read/write rates (millions of ops per second)
- Efficient scans over all the data or interesting subsets

❖ Temporal perspective:

- Examine data changes over time,
e.g.: contents of a web page over multiple crawls.

Characteristics of Bigtable

- ❖ Distributed, multi-level, sparse map:
 - Row key – column – timestamp
- ❖ Fault-tolerant, persistent
- ❖ Performance:
 - Fast access to very large amounts of data (structured, semi-structured, unstructured)
 - High read and write throughput
- ❖ Scalability:
 - Dynamically add / remove cluster nodes
 - Scale seamlessly from thousands to millions of reads/writes per second
 - Automatic replication

Building Blocks (from Google Infrastructure)

- ❖ **Google File System (GFS):** to store log and data files
- ❖ **SSTable (Sorted Strings Table):**
 - immutable file format to store log and data files
 - Key/value string pairs sorted by key
 - consists of sequence of blocks (64kB default) + block index (loaded to memory when SSTable is opened).
- ❖ **Chubby:**
 - highly available, persistent distributed lock service
 - 5 active replicas with 1 elected master
 - entry point to Bigtable data

Contents

❖ Column Stores:

- Motivation
- Explicit vs. Implicit IDs
- Compression
- Late Materialization
- Updates

❖ Column Family Stores (Wide Column Stores):

- Motivation
- **Data Model**

Rows

- ❖ Data is stored in **lexicographic order by key**
- ❖ Key can be arbitrary string ($\leq 64\text{kB}$, usually much less)
- ❖ Rows with consecutive keys are grouped into **tablets** (= the unit of distribution and load balancing).
 - "neighbouring" rows on 1 or small number of nodes
 - applications can choose key to achieve **good locality** (e.g., store web pages with reversed url as key).
- ❖ Each tablet consists of one or more **SSTables**.
- ❖ Access to data in a row is **atomic**.

Columns


❖ Columns are grouped into column families:

- Two level name structure: *family:optional_qualifier*
- Small number of column families recommended (100s)
- Unbounded number of columns in a family

❖ Column family:

- Unit of access control
- Must be created explicitly before data can be stored.

Example:



Timestamps

- ❖ Cells in a table can contain multiple versions of data
 - Can be assigned implicitly by Bigtable ("real time" microseconds) or explicitly by the application.
 - Cell versions are stored in decreasing timestamp order.
- ❖ Lookup options:
 - "Return most recent K values" or
 - "Return all values in a given timestamp range" or all.
- ❖ Garbage collection options (per column family):
 - "Only retain most recent K values"
 - "Retain values until they are older than K seconds".

Data Access via API

❖ Schema:

- Create/delete tables or column families
- Change metadata (such as access rights)

❖ Read access:

- Single-row reads (based on row key)
- Scans (range of rows, multiple ranges of rows)

❖ Write access:

- Simple write, batch write (ideally contiguous rows)
- Modifications (e.g., increment integers, append data)

❖ Read/write data in Bigtable with MapReduce

Lessons Learned (according to Google)

❖ Prepare for many types of failure:

- not just single node failure or network partition
- also e.g. memory/network corruption, huge clock skew, hung machines, bugs in other systems, etc.

❖ Delay adding new features until their usage is clear:

- Originally, general purpose transactions were planned;
- ultimately, mostly single-row transactions were needed.

❖ Importance of proper system-level monitoring

- Bigtable itself, but also client processes using Bigtable

❖ "Most important" lesson: simple design

Further Popular Systems

- ❖ **HBase**: Bigtable-like open source implementation
 - PC/EC according to PACELC categorization
 - master-based
 - access only via APIs (no in-built query language)
 - built on top of Hadoop and HDFS:
good integration with Hive, Spark, etc.
- ❖ **Cassandra**: quite different from Bigtable
 - PA/EL according to PACELC categorization
 - masterless, data model different from Bigtable
 - Cassandra Query Language (CQL): SQL-like

Summary

- ❖ **Column stores:** basic ideas, motivation
- ❖ **Basic principles of column stores:**
 - Explicit vs. implicit ids
 - Compression: motivation, methods
 - Materialization: early vs. late materialization
 - Updating columns: challenges, solutions
- ❖ **Wide column stores (column family stores):**
 - Motivation, requirements, characteristics
 - Data model and data access
 - Common systems: Bigtable, Cassandra, HBase